

Insurance Fraud Project

May 24, 2023

Objective of this project is to find the best ML Model that can predict Car Insurance Frauds. This dataset was found on Kaggle.com.

First I'll explore the dataset in order to show its characteristic with EDA. Second, I'll preprocess the data, then I'll train and test 6 ML/DL models.

0.1 Importing necessary libraries and the dataset

```
[1]: # Importing necessary libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
from imblearn.under_sampling import RandomUnderSampler
from tabulate import tabulate
from sklearn.metrics import roc_auc_score, confusion_matrix
from matplotlib.colors import ListedColormap

# Reading the dataset
df = pd.read_csv("C:/Users/Utente/Desktop/Progetto/fraud_oracle.csv")

[2]: df.head()
```

```
[2]:   Month  WeekOfMonth  DayOfWeek  Make  AccidentArea  DayOfWeekClaimed  \
0   Dec             5  Wednesday  Honda         Urban         Tuesday
1   Jan             3  Wednesday  Honda         Urban          Monday
2   Oct             5    Friday   Honda         Urban        Thursday
3   Jun             2  Saturday  Toyota         Rural          Friday
4   Jan             5    Monday   Honda         Urban         Tuesday
```

	MonthClaimed	WeekOfMonthClaimed	Sex	MaritalStatus	...	AgeOfVehicle	\
0	Jan	1	Female	Single	...	3 years	
1	Jan	4	Male	Single	...	6 years	
2	Nov	2	Male	Married	...	7 years	
3	Jul	1	Male	Married	...	more than 7	
4	Feb	2	Female	Single	...	5 years	

	AgeOfPolicyHolder	PoliceReportFiled	WitnessPresent	AgentType	\
0	26 to 30	No	No	External	
1	31 to 35	Yes	No	External	
2	41 to 50	No	No	External	
3	51 to 65	Yes	No	External	
4	31 to 35	No	No	External	

	NumberOfSuppliments	AddressChange_Claim	NumberOfCars	Year	BasePolicy
0	none	1 year	3 to 4	1994	Liability
1	none	no change	1 vehicle	1994	Collision
2	none	no change	1 vehicle	1994	Collision
3	more than 5	no change	1 vehicle	1994	Liability
4	none	no change	1 vehicle	1994	Collision

[5 rows x 33 columns]

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15420 entries, 0 to 15419
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Month                 15420 non-null  object
1   WeekOfMonth           15420 non-null  int64
2   DayOfWeek             15420 non-null  object
3   Make                  15420 non-null  object
4   AccidentArea          15420 non-null  object
5   DayOfWeekClaimed      15420 non-null  object
6   MonthClaimed          15420 non-null  object
7   WeekOfMonthClaimed    15420 non-null  int64
8   Sex                   15420 non-null  object
9   MaritalStatus         15420 non-null  object
10  Age                   15420 non-null  int64
11  Fault                 15420 non-null  object
12  PolicyType            15420 non-null  object
13  VehicleCategory       15420 non-null  object
14  VehiclePrice          15420 non-null  object
15  FraudFound_P         15420 non-null  int64
16  PolicyNumber          15420 non-null  int64
17  RepNumber             15420 non-null  int64
```

```

18 Deductible          15420 non-null int64
19 DriverRating        15420 non-null int64
20 Days_Policy_Accident 15420 non-null object
21 Days_Policy_Claim    15420 non-null object
22 PastNumberOfClaims   15420 non-null object
23 AgeOfVehicle         15420 non-null object
24 AgeOfPolicyHolder    15420 non-null object
25 PoliceReportFiled    15420 non-null object
26 WitnessPresent      15420 non-null object
27 AgentType           15420 non-null object
28 NumberOfSupplements  15420 non-null object
29 AddressChange_Claim  15420 non-null object
30 NumberOfCars         15420 non-null object
31 Year                15420 non-null int64
32 BasePolicy          15420 non-null object
dtypes: int64(9), object(24)
memory usage: 3.9+ MB

```

```
[4]: df.describe()
```

```

[4]:      WeekOfMonth  WeekOfMonthClaimed      Age  FraudFound_P \
count  15420.000000      15420.000000  15420.000000  15420.000000
mean      2.788586      2.693969    39.855707    0.059857
std      1.287585      1.259115    13.492377    0.237230
min      1.000000      1.000000     0.000000    0.000000
25%      2.000000      2.000000    31.000000    0.000000
50%      3.000000      3.000000    38.000000    0.000000
75%      4.000000      4.000000    48.000000    0.000000
max      5.000000      5.000000    80.000000    1.000000

      PolicyNumber  RepNumber  Deductible  DriverRating      Year
count  15420.000000  15420.000000  15420.000000  15420.000000  15420.000000
mean    7710.500000    8.483268   407.704280    2.487808  1994.866472
std   4451.514911    4.599948   43.950998    1.119453    0.803313
min      1.000000    1.000000   300.000000    1.000000  1994.000000
25%   3855.750000    5.000000   400.000000    1.000000  1994.000000
50%   7710.500000    8.000000   400.000000    2.000000  1995.000000
75%  11565.250000   12.000000   400.000000    3.000000  1996.000000
max  15420.000000   16.000000   700.000000    4.000000  1996.000000

```

0.2 Data Preprocessing and EDA

```


[5]: # Function to check missing values
def check_na(data):
    na = data.isna().sum()
    return na

```


```
check_na(df)
```

```
[5]: Month                0
     WeekOfMonth          0
     DayOfWeek            0
     Make                 0
     AccidentArea         0
     DayOfWeekClaimed     0
     MonthClaimed         0
     WeekOfMonthClaimed   0
     Sex                  0
     MaritalStatus        0
     Age                  0
     Fault                0
     PolicyType           0
     VehicleCategory      0
     VehiclePrice         0
     FraudFound_P         0
     PolicyNumber         0
     RepNumber            0
     Deductible           0
     DriverRating         0
     Days_Policy_Accident 0
     Days_Policy_Claim    0
     PastNumberOfClaims   0
     AgeOfVehicle         0
     AgeOfPolicyHolder    0
     PoliceReportFiled    0
     WitnessPresent       0
     AgentType            0
     NumberOfSupplements  0
     AddressChange_Claim  0
     NumberOfCars         0
     Year                 0
     BasePolicy           0
     dtype: int64
```

```
[6]: # Create the countplot for Age Variable
plt.figure(figsize=(10, 6))
ax = sns.countplot(x='Age', data=df)
ax.tick_params(axis='x', labelsize=6)
plt.xlabel('Age')
plt.ylabel('Count')
plt.title('Age Distribution')
plt.show()
```




```
[7]: plt.figure(figsize=(10, 6))
df.groupby('Age')['FraudFound_P'].count().plot(kind='bar')
plt.title('Count of Frauds by Age')
plt.xlabel('Age')
plt.ylabel('Count of Fraud')
plt.show()
```



```
[8]: # Calculate the percentage of frauds for each age category
fraud_percentage = df[df['FraudFound_P'] == 1].groupby('Age')['FraudFound_P'].
    ↪count() / df.groupby('Age')['FraudFound_P'].count()

# Create the bar plot
plt.figure(figsize=(10, 6))
ax = sns.barplot(x=fraud_percentage.index, y=fraud_percentage.values,
    ↪color='blue')
ax.tick_params(axis='x', labelsz=8)
plt.xlabel('Age')
plt.ylabel('Percentage of Frauds')
plt.title('Percentage of Frauds by Age')

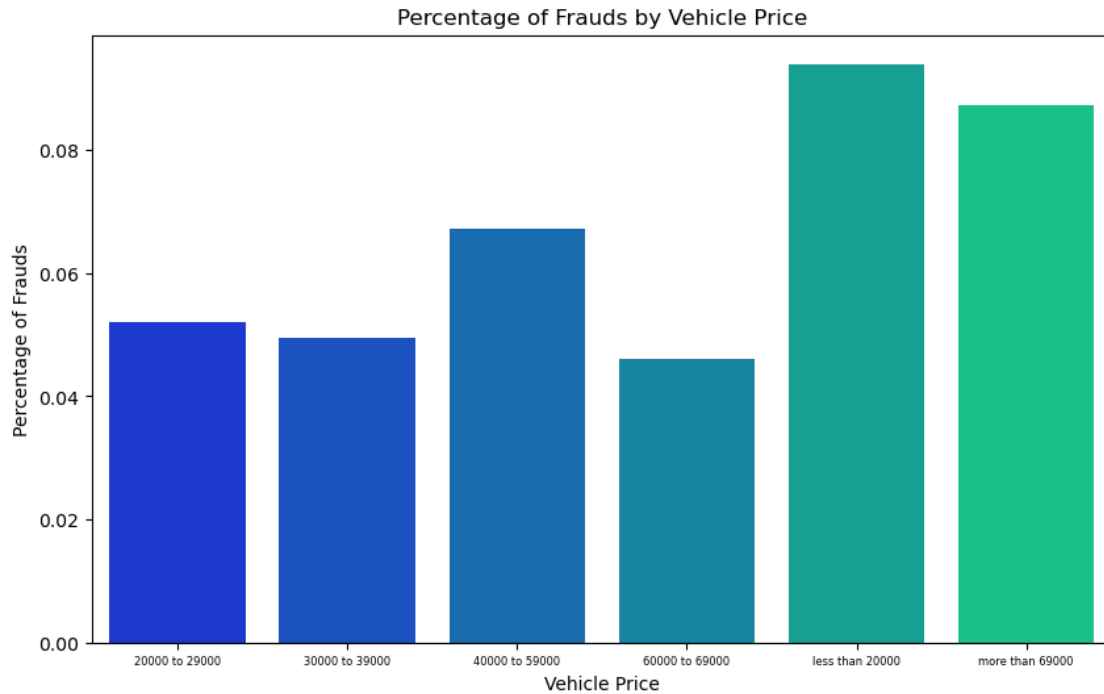
plt.show()
```



```
[9]: # Calculate the percentage of frauds for each age category
fraud_percentage = df[df['FraudFound_P'] == 1].
    ↳groupby('VehiclePrice')['FraudFound_P'].count() / df.
    ↳groupby('VehiclePrice')['FraudFound_P'].count()

# Create the bar plot
plt.figure(figsize=(10, 6))
pal = sns.color_palette("winter", len(fraud_percentage))
ax = sns.barplot(x=fraud_percentage.index, y=fraud_percentage.values,
    ↳palette=pal)
ax.tick_params(axis='x', labels=6)
plt.xlabel('Vehicle Price')
plt.ylabel('Percentage of Frauds')
plt.title('Percentage of Frauds by Vehicle Price')

# Display the plot
plt.show()
```




```
[10]: # Calculate the percentage of frauds for each sex category
fraud_percentage = df[df['FraudFound_P'] == 1]['Sex'].
    ↪ value_counts(normalize=True) * 100

# Create the pie plot
plt.figure(figsize=(8, 6))
colors = ['skyblue', 'lightcoral']
plt.pie(fraud_percentage, labels=fraud_percentage.index, autopct='%1.1f%%',
    ↪ colors=colors)
plt.title('Percentage of Frauds by Sex')

# Display the plot
plt.show()
```


Percentage of Frauds by Sex



```
[11]: # Creating dummy variables for categorical columns
columns_dummies = ['AccidentArea', 'AgeOfPolicyHolder', 'BasePolicy', 'Make',
                    'MaritalStatus', 'Sex', 'Days_Policy_Claim',
                    ↪ 'PastNumberOfClaims',
                    'VehiclePrice', 'AgeOfVehicle', 'AddressChange_Claim',
                    ↪ 'PoliceReportFiled',
                    'WitnessPresent', 'AgentType', 'NumberOfSupplements',
                    'NumberOfCars', 'VehicleCategory', 'Fault', 'PolicyType']

df = pd.concat([df.drop(columns_dummies, axis=1),
                pd.get_dummies(df[columns_dummies], dtype=float)], axis=1)

# Mapping categorical values to numerical values
day_mapping = {
    'Monday': 1,
    'Tuesday': 2,
    'Wednesday': 3,
```

```

        'Thursday': 4,
        'Friday': 5,
        'Saturday': 6,
        'Sunday': 7
    }

df['DayOfWeek'] = df['DayOfWeek'].map(day_mapping)
df['DayOfWeekClaimed'] = df['DayOfWeekClaimed'].map(day_mapping)
print(df['DayOfWeek'])

month_mapping = {
    'Jan': 1,
    'Feb': 2,
    'Mar': 3,
    'Apr': 4,
    'May': 5,
    'Jun': 6,
    'Jul': 7,
    'Aug': 8,
    'Sep': 9,
    'Oct': 10,
    'Nov': 11,
    'Dec': 12
}

df['Month'] = df['Month'].map(month_mapping)
df['MonthClaimed'] = df['MonthClaimed'].map(month_mapping)

# Checking missing values again
def check_na(data):
    na = data.isna().sum()
    return na

check_na(df)

# Dropping rows with missing values
df = df.dropna()
check_na(df)
df.head()

```

```

0      3
1      3
2      5
3      6
4      1
..
15415  5
15416  4

```

```
15417    4
15418    1
15419    3
```

Name: DayOfWeek, Length: 15420, dtype: int64

```
[11]:   Month  WeekOfMonth  DayOfWeek  DayOfWeekClaimed  MonthClaimed  \
0      12           5           3             2.0           1.0
1       1           3           3             1.0           1.0
2      10           5           5             4.0          11.0
3       6           2           6             5.0           7.0
4       1           5           1             2.0           2.0

      WeekOfMonthClaimed  Age  FraudFound_P  PolicyNumber  RepNumber  ...  \
0                      1   21             0             1          12  ...
1                      4   34             0             2          15  ...
2                      2   47             0             3           7  ...
3                      1   65             0             4           4  ...
4                      2   27             0             5           3  ...

      Fault_Third Party  PolicyType_Sedan - All Perils  \
0                    0.0                        0.0
1                    0.0                        0.0
2                    0.0                        0.0
3                    1.0                        0.0
4                    1.0                        0.0

      PolicyType_Sedan - Collision  PolicyType_Sedan - Liability  \
0                        0.0                        0.0
1                        0.0                        0.0
2                        0.0                        0.0
3                        0.0                        1.0
4                        0.0                        0.0

      PolicyType_Sport - All Perils  PolicyType_Sport - Collision  \
0                        0.0                        0.0
1                        0.0                        1.0
2                        0.0                        1.0
3                        0.0                        0.0
4                        0.0                        1.0

      PolicyType_Sport - Liability  PolicyType_Utility - All Perils  \
0                        1.0                        0.0
1                        0.0                        0.0
2                        0.0                        0.0
3                        0.0                        0.0
4                        0.0                        0.0
```

	PolicyType_Utility - Collision	PolicyType_Utility - Liability
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

[5 rows x 109 columns]

```
[12]: # Plotting the count of fraud and non-fraud transactions
ax = sns.countplot(x='FraudFound_P', data=df)
plt.xlabel('Fraud')
plt.ylabel('Count')
plt.title('Number of non-fraud and fraud transactions')

# Adding count numbers inside each bin
for p in ax.patches:
    height = p.get_height()
    ax.annotate(f'{height}', (p.get_x() + p.get_width() / 2, height),
                ha='center', va='bottom', fontweight='bold', color='black')

# Displaying the plot
plt.show()
```



0.3 Rebalancing the Dataset

```
[13]: # Calculating the ratio of fraud transactions
unbalance = df.FraudFound_P[df['FraudFound_P']==1].count()/df['FraudFound_P'].
        ↪count()
print(unbalance)
```

0.0598612101952137

```
[14]: X = df.drop('FraudFound_P', axis=1)
y = df['FraudFound_P']

rus = RandomUnderSampler()

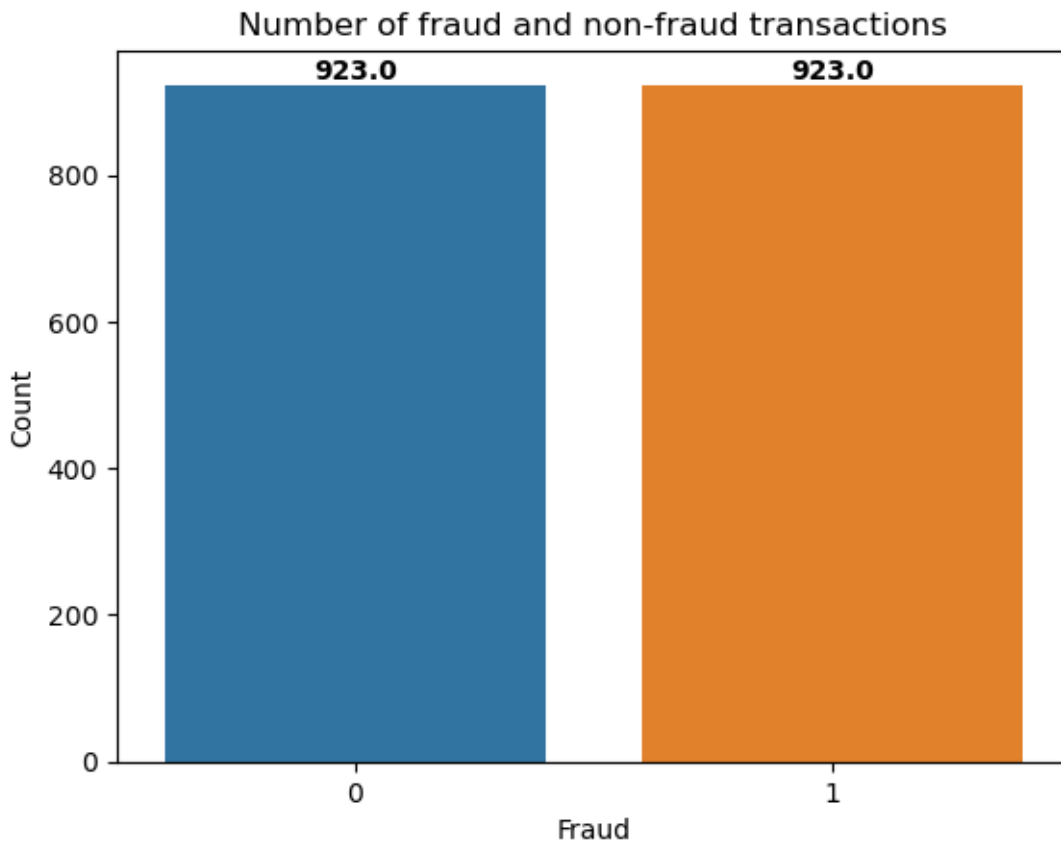
X_resampled, y_resampled = rus.fit_resample(X, y)
df = pd.concat([X_resampled, y_resampled], axis=1)
ax = sns.countplot(x='FraudFound_P', data=df)

# add labels and title
plt.xlabel('Fraud')
```

```
plt.ylabel('Count')
plt.title('Number of fraud and non-fraud transactions')

# add count numbers inside each bin
for p in ax.patches:
    height = p.get_height()
    ax.annotate(f'{height}', (p.get_x() + p.get_width() / 2, height),
                ha='center', va='bottom', fontweight='bold', color='black')

# display the plot
plt.show()
```



0.4 Splitting, Scaling and Encoding Data

```
[15]: # Splitting the data into training and testing sets
X = df.drop('FraudFound_P', axis=1)
y = df['FraudFound_P']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```

# Scaling numerical features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train.select_dtypes(include=['float64',
↪ 'int64']))
X_test_scaled = scaler.transform(X_test.select_dtypes(include=['float64',
↪ 'int64']))

# Encoding target variable
encoder = LabelEncoder()
y_train_enc = encoder.fit_transform(y_train)
y_test_enc = encoder.transform(y_test)

```

0.5 Training and Evaluating Models

```

[16]: # Define the hyperparameter grids for each model
lr_param_grid = {'C': [0.1, 1, 10], 'penalty': ['l2']}
rf_param_grid = {'n_estimators': [100, 200, 300], 'max_depth': [None, 5, 10]}
dt_param_grid = {'max_depth': [None, 5, 10], 'min_samples_split': [2, 5, 10]}
nn_param_grid = {'hidden_layer_sizes': [(100,), (100, 50)], 'alpha': [0.0001, 0.
↪ 0.001, 0.01]}
gb_param_grid = {'n_estimators': [100, 200, 300], 'learning_rate': [0.1, 0.05,
↪ 0.01]}
svc_param_grid = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}

# Define the models
models = [
    (LogisticRegression(max_iter=1000), lr_param_grid),
    (RandomForestClassifier(), rf_param_grid),
    (DecisionTreeClassifier(), dt_param_grid),
    (MLPClassifier(max_iter=1000), nn_param_grid),
    (SVC(), svc_param_grid),
    (GradientBoostingClassifier(), gb_param_grid)
]

```

```

[17]: # Perform model selection and evaluation
results = []
metrics = ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']

for model, param_grid in models:
    scoring = 'roc_auc'

    grid_search = GridSearchCV(model, param_grid, cv=5, scoring=scoring,
↪ n_jobs=-1)
    grid_search.fit(X_train_scaled, y_train_enc)

    best_model = grid_search.best_estimator_
    best_model.fit(X_train_scaled, y_train_enc) # Fit the best model

```

```

y_pred = best_model.predict(X_test_scaled)

row = []
for metric in metrics:
    score = eval(metric + '_score')(y_test_enc, y_pred)
    row.append(score)

results.append(row)

df_results = pd.DataFrame(results, columns=metrics, index=[best_model.__class__.__name__ for best_model, _ in models])

# Display the results in a table
print(tabulate(df_results, headers='keys', tablefmt='psql'))

```

```

+-----+-----+-----+-----+-----+
|          | accuracy | precision | recall | f1 |
|-----|
| LogisticRegression | 0.77027 | 0.726496 | 0.890052 | 0.8 |
| RandomForestClassifier | 0.802703 | 0.741803 | 0.947644 | 0.832184 |
| DecisionTreeClassifier | 0.745946 | 0.715556 | 0.842932 | 0.774038 |
| MLPClassifier | 0.708108 | 0.724324 | 0.701571 | 0.712766 |
| SVC | 0.77027 | 0.71371 | 0.926702 | 0.806378 |
| GradientBoostingClassifier | 0.894595 | 0.85514 | 0.958115 | 0.903704 |

```

```

[18]: # Plotting the evaluation metrics for different models
scores = np.array(results)
colors = ListedColormap([
    '#00008B', '#002180', '#004575', '#00686A',
    '#008B5F', '#00AE54', '#00D149', '#00F53E',
    '#00FF00'
])

fig, axs = plt.subplots(5, 1, figsize=(8, 20))
metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score', 'ROC AUC']

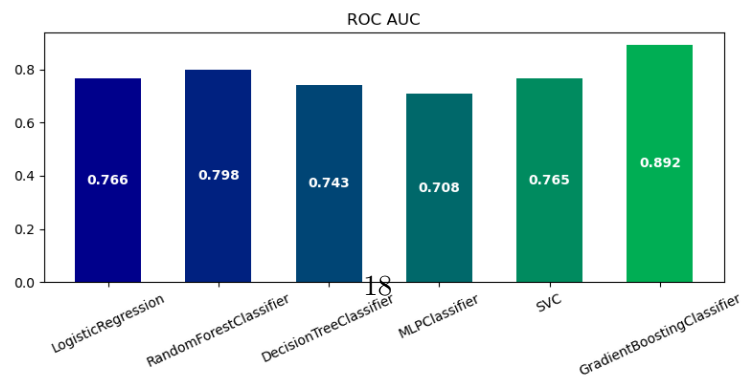
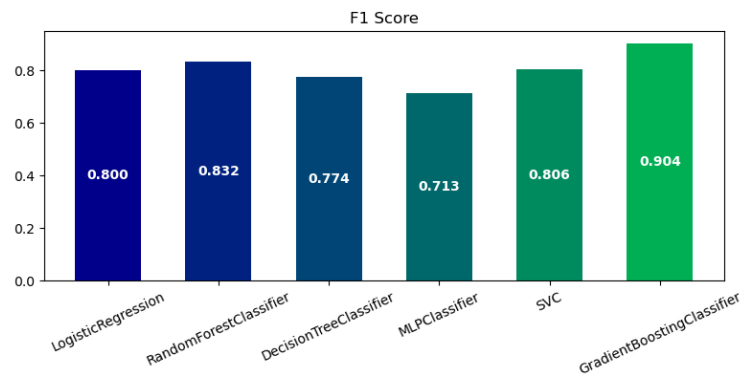
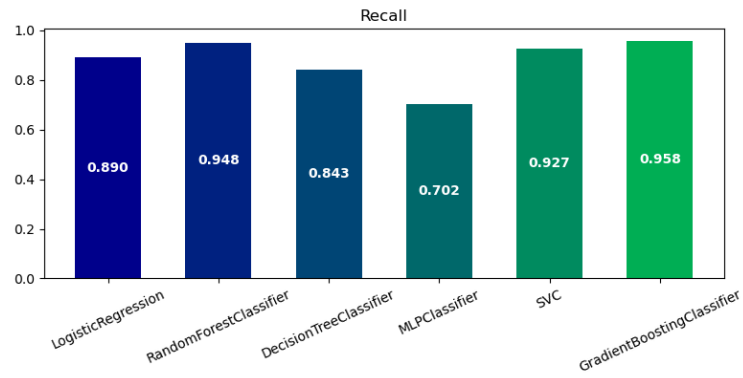
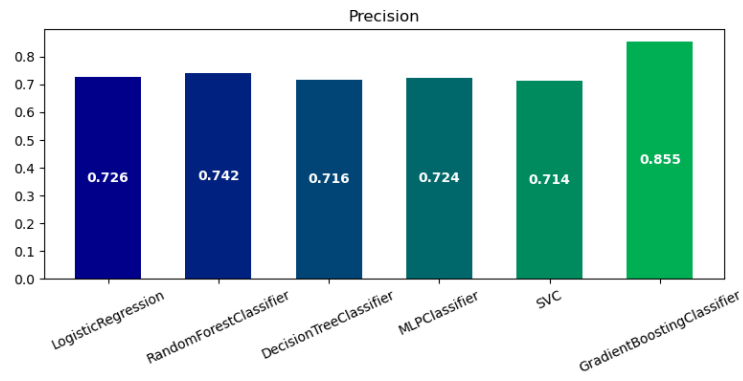
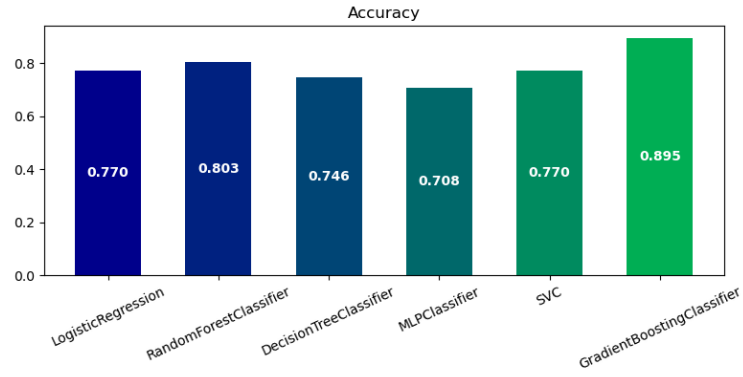
```



```

for i, metric in enumerate(metrics):
    bars = axs[i].bar([best_model.__class__.__name__ for best_model, _ in
↳models], scores[:, i], color=colors(range(len(models))), width=0.6)
    axs[i].set_title(metric)
    axs[i].tick_params(axis='x', rotation=25)
    for bar in bars:
        height = bar.get_height()
        axs[i].text(bar.get_x() + bar.get_width() / 2, height * 0.5, f'{height:.
↳3f}', ha='center', va='center', color='white', fontweight='bold')
plt.tight_layout()
plt.show()

```



```

[19]: # Generating and displaying confusion matrices for different models
num_models = len(models)
num_cols = 2 # Number of columns to display

# Calculate the number of rows needed to display all the models
num_rows = (num_models + num_cols - 1) // num_cols

# Set the figure size based on the number of rows and columns
fig, axes = plt.subplots(num_rows, num_cols, figsize=(8, 6*num_rows))

# Flatten the axes array for easy iteration
axes = axes.flatten()

for i, (best_model, _) in enumerate(models):
    best_model.fit(X_train_scaled, y_train_enc)
    y_pred = best_model.predict(X_test_scaled)
    cm = confusion_matrix(y_test_enc, y_pred)

    # Set the current subplot
    ax = axes[i]

    # Plot the confusion matrix
    im = ax.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
    ax.set_xticks([0, 1])
    ax.set_yticks([0, 1])
    ax.set_xticklabels(['Not Fraud', 'Fraud'])
    ax.set_yticklabels(['Not Fraud', 'Fraud'])
    ax.set_xlabel('Predicted label')
    ax.set_ylabel('True label')
    ax.set_title(f'Confusion Matrix for {best_model.__class__.__name__}')







    for i in range(2):
        for j in range(2):
            ax.text(j, i, str(cm[i, j]), ha='center', va='center',
                    color='black', fontweight='bold')

# Remove empty subplots if there are any
if num_models < len(axes):
    for i in range(num_models, len(axes)):
        fig.delaxes(axes[i])

# Adjust the spacing between subplots
plt.tight_layout()

```

```
# Display the plot  
plt.show()
```



As we can see the best model is the Gradient Boosting Classifier with this parameters:

```
[20]: best_model_params = grid_search.best_params_  
print("Best Model Parameters:")  
print(best_model_params)
```

Best Model Parameters:

```
{'learning_rate': 0.1, 'n_estimators': 300}
```